

Shortest-Path Problem


EECS 233

Depth-first Traversal

- Visit a vertex, then make recursive calls on all of its yet-to-be-visited neighbors:


depthFirstTrav(v)
myDepthFirstTrav(v, NULL)

myDepthFirstTrav(node, parent)
visit node and mark it as visited
node.parent = parent
for each vertex w in node's neighbors
if (w has not been visited)
myDepthFirstTrav(w, node)



```
public class Graph {
    class Vertex {
        private String id;
        private linkedList <Edge> edges;
        // adjacency list
        private boolean encountered;
        private boolean done;
        private Vertex parent;
        private double cost;
        ...
    }
    class Edge {
        private int endNode;
        private double cost;
        ...
    }
    private Vertex[] vertices;
    private int numVertices;
    private int maxNum;
    ...
}
```

Breadth-First Traversal



Order: Cincinnati, Columbus, Toledo, Indianapolis, Pittsburgh, Cleveland, Detroit, Chicago, Buffalo

- Starting from Cincinnati, what would be the order of visits?
- **breadth-first:**
 - visit a vertex
 - visit all of its neighbors
 - visit all unvisited vertices 2 edges away
 - visit all unvisited vertices 3 edges away, etc.

The Shortest-Path Problem


- What is the “shortest” path from one vertex to another – i.e., the one with the minimal total cost
 - Travel: lowest mileage or fastest trip
 - Internet: forwarding traffic along the “best” router paths.
- Given a source, find the shortest path to a destination
- Given a source, find the shortest path to many (all) destinations

First, A Special Case

- A solution for an *unweighted* graph:


First, A Special Case

- A solution for an *unweighted* graph:



First, A Special Case

- A solution for an *unweighted* graph:
 - Start a breadth-first traversal from the origin, *x*
 - Stop the traversal when you reach the target, *y*
 - The path from *x* to *y* in the resulting (possibly partial) spanning tree is *a* shortest path
- A breadth-first traversal works for an unweighted graph because
 - the shortest path is simply one with the fewest edges
 - a breadth-first traversal visits nodes in order according to the number of edges they are from the origin.



Method for the Special Case (Pseudo-code)

```
shortestPath_bfTrav(source, desti)
    mark source as encountered
    source.parent = null
    source.cost = 0;
    create a new queue q
    q.insert(source);


    while (!q.isEmpty())
        v = q.remove()
        if (v==desti) break;
        for each vertex w in v's neighbors
            if w is not encountered
                mark w as encountered
                w.parent = v;
                w.cost = v.cost + 1;
                q.insert(w);
```


Shortest paths for Weighted Graphs: Dijkstra's Algorithm

- Finds the shortest paths from vertex v to *all other vertices* that can be reached from v .
 - The single-source-all-destinations problem
- Assumptions:
 - Graph is connected.
 - Edge cost is non-negative
 - More hops = higher cost
- Basic idea:
 - maintain estimates of the shortest paths from v to every vertex (along with their associated costs)
 - gradually refine these estimates as we traverse the graph


Dijkstra's Algorithm

- We say that a vertex w is **finalized** if we are certain we have found the shortest path from v to w .
- We repeatedly do the following:
 - find the not-yet-finalized vertex w with the lowest cost estimate
 - mark w as finalized (shown as a filled circle below)
 - examine each not-yet-finalized neighbor x of w to see if there is a shorter path to x that passes through w ; if there is, update the shortest-path estimate for x

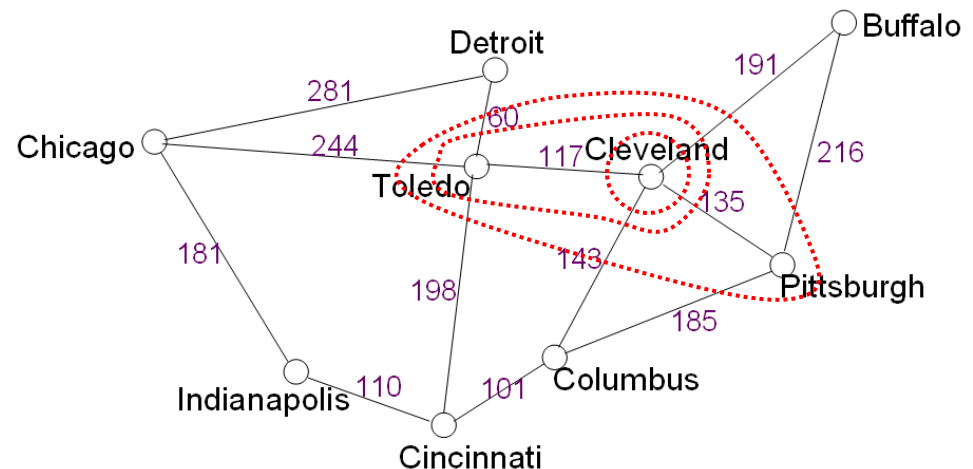


Intuition Behind Dijkstra's

- For an unweighted graph, the breadth-first traversal method expands the search circle by one hop every round

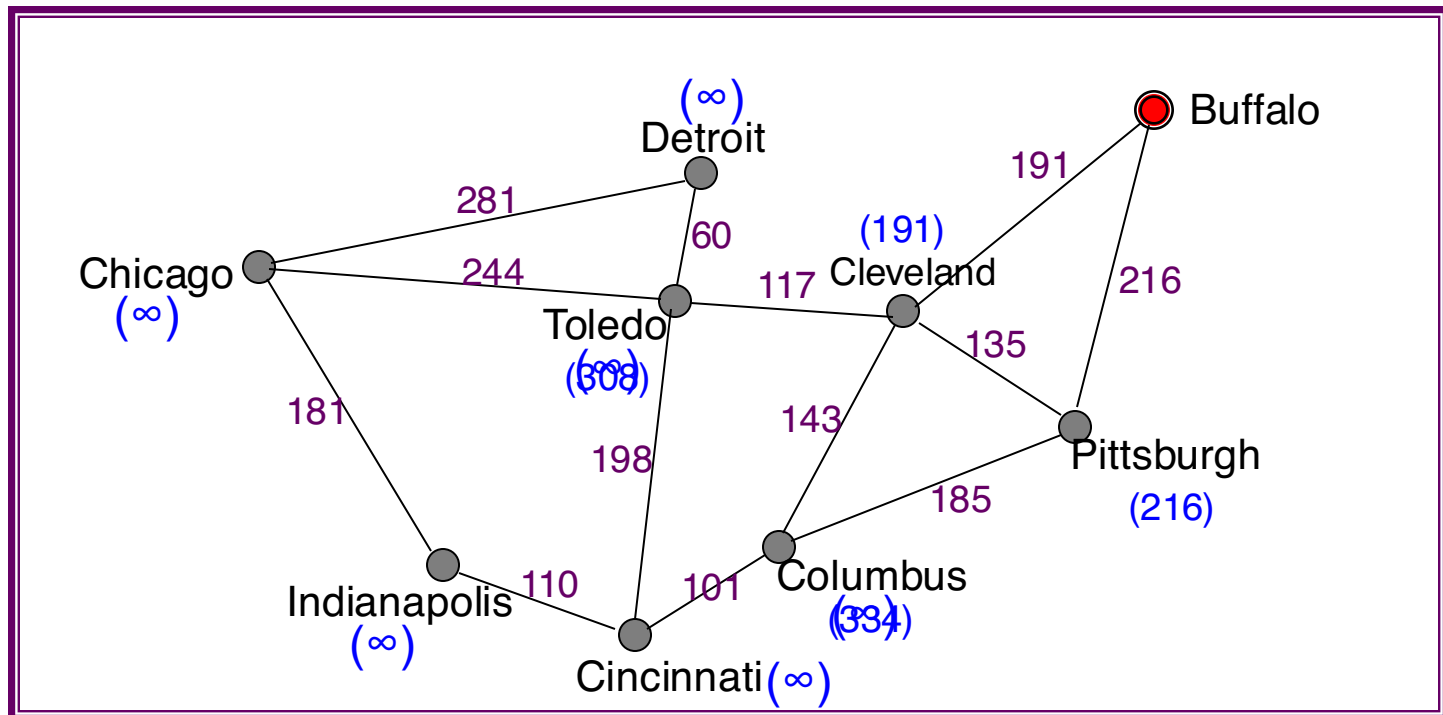


- For a weighted graph, Dijkstra's method expands the search circle **by one lowest-cost node** every round



Example: Shortest Path from Buffalo To Toledo

- Distance to Buffalo is 0, so it is finalized
- Initially set distance estimates:
 - To all neighbors - edge weights
 - To other cities - unknown, using *infinity*
- Which will be finalized?



Dijkstra's Algorithm

Notation:

$c(x,y)$: link cost from node x to y ;

∞ if not direct neighbors

$D(v)$: current path cost estimate from source to dest. v

$p(v)$: predecessor node along path from source to v

N : set of nodes whose least cost path definitively known

u : starting node

1 Initialization:

2 $N = \{u\}$

3 for all nodes v

4 if v adjacent to u then

5 $D(v) = c(u,v)$

6 $p(v) = u$

7 else $D(v) = \infty$

8 Loop

9 find w not in N with smallest $D(w)$

10 add w to N

11 update $D(v)$ for all v adjacent to w and not in N :


12 $D(v) = \min(D(v), D(w) + c(w,v))$

Update $p(v)$

13 /* new cost to v is either old cost to v or known



14 shortest path cost to w plus cost from w to v */

15 until all nodes are in N



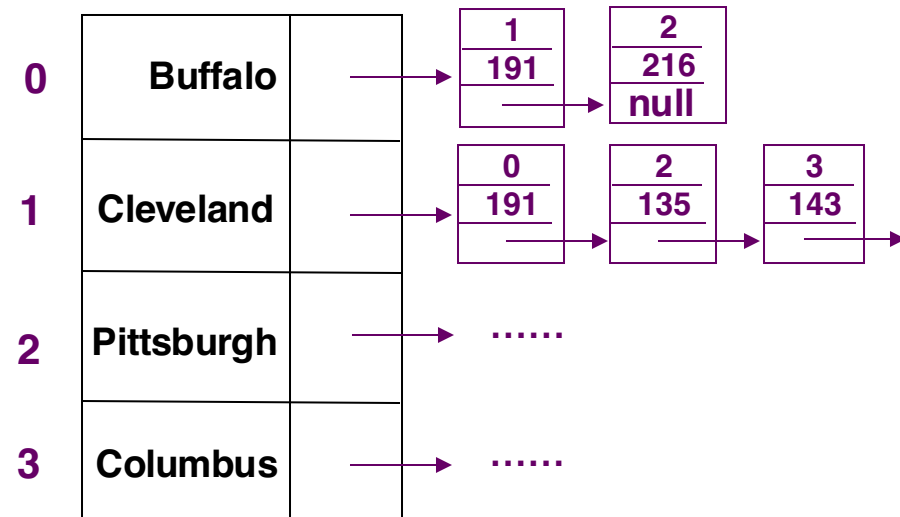
Dijkstra's algorithm: example

Step	N	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	2, u	5, u	1, u	∞	∞



Implementation Issues


- Data structure: adjacency lists



- How to find the not-yet-finalized vertex with the minimal cost?
What data structure can be used to make it more efficient?

Dijkstra's Algorithm: Complexity

```
1 Initialization:
2   N = {u}
3   for all nodes v           // V iterations
4     if v adjacent to u      // Check of the total of up to E edges
5       then D(v) = c(u,v)
6     else D(v) = ∞
7                               // O(V) to build the heap
8 Loop           // V iterations
9   find w not in N with smallest D(w)           // logV to reorg
10  add w to N                                   the heap
11  update D(v) for all v adjacent to w and not in N :
12    D(v) = min( D(v), D(w) + c(w,v) )           // A step per every link (E total); logV
13    /* new cost to v is either old cost to v or known           per step to reorganize the heap
14    shortest path cost to w plus cost from w to v */
15 until all nodes in N
```




Overall complexity is $O(V + E + V \log V + E \log V) = O(E \log V)$

All-Pairs Shortest-Path Problem

- Dijkstra's algorithm solves the problem of **single**-source all-destination shortest paths
- It can be used to solve the all-source **all**-destination shortest path problem too
 - Run Dijkstra's V times, once for each vertex
 - Running time: $O(VE \log V)$

Shortest Paths



A Weighted Directed Graph


(b)

	0	1	2	3	4
0	∞	8	∞	9	4
1	∞	∞	1	∞	∞
2	∞	2	∞	3	∞
3	∞	∞	2	∞	7
4	∞	∞	1	∞	∞

Its Adjacency Matrix

Dijkstra's Shortest-Path Algorithm – Trace

Step	v	vertexSet	weight				
			[0]	[1]	[2]	[3]	[4]
1	–	0	0	8	∞	9	4
2	4	0, 4	0	8	5	9	4
3	2	0, 4, 2	0	7	5	8	4
4	1	0, 4, 2, 1	0	7	5	8	4
5	3	0, 4, 2, 1, 3	0	7	5	8	4



(b)

	0	1	2	3	4
0	∞	8	∞	9	4
1	∞	∞	1	∞	∞
2	∞	2	∞	3	∞
3	∞	∞	2	∞	7
4	∞	∞	1	∞	∞